



# Docker Toolbox

Hands-On Workshop



# Agenda

---

- LESSON 1 - PART 1 Intro to Docker (15 min)
  - Docker, Cgroups, Kernel Namespaces, and Docker Images
- LESSON 1 - PART 2 – Toolbox Overview (15 min)
  - Docker Hub, Trusted build system, Weave, Flannel, etcd, CoreOS, Magnum, Kubernetes, and how Docker works with OpenStack
- Hands-on LAB 1 - Using docker-machine (15 min)
  - Spin up Docker containers on Rackspace Cloud, Switch between docker hosts.
  - How to run ad-hoc containers on the fly for experimentation, How to get a shell on a Docker host, and a Docker container without running sshd in the container.
  - How to move a container from one host to another.
- LESSON 2 – Writing Dockerfiles (15 min)
  - Dockerfile Directives
  - Demonstration of copying files into containers using bind mounts, and copying files out of containers using “docker cp”.
  - Why Dockerfiles are better than customized images made manually and stored with commit.
- Hands-on LAB 2 - Writing Dockerfiles (15 min)
  - Saving a Dockerfile in a source repo, and using a build script for configuration injection.
- LESSON 3 - Linking and Networking containers (15 min)
  - Demonstration of how to use links to communicate between containers on the same host. Using weave to communicate between containers on separate hosts.
- Hands on LAB 3 - Linking containers (15 min)
  - Container linking on the same host, and mesh networking between containers on different hosts.

# 1.1


## **Lesson One – Part One**

### **Introduction to Docker**

# Container

---

- Combines several things
  - Linux Cgroups
  - Kernel Namespaces
  - Docker Image
  - Has a lifecycle



# Linux Cgroups

---

- Kernel Feature
- Groups of processes
- Control resource allocations
  - CPU
  - Memory
  - Disk
  - I/O
- May be nested



# Linux Kernel Namespaces

---


- Kernel Feature
- Restrict your view of the system
  - Mounts (CLONE\_NEWNS)
  - UTS (CLONE\_NEWUTS)
    - `uname ( )` output
  - IPC (CLONE\_NEWIPC)
  - PID (CLONE\_NEWPID)
  - Networks (CLONE\_NEWNET)
  - User (CLONE\_NEWUSER)
    - Not supported in Docker yet
    - Has privileged/unprivileged modes today
- May be nested



## Docker Image

---


- NOT A FILESYSTEM
- NOT A VHD
- Basically a tar file
- Has a hierarchy
  - Arbitrary depth
- Fits into the Docker Registry



# Docker Registry

---

- Git Repo Semantics
  - Pull
  - Push
  - Commit
  - Hierarchy






# Container

---

- Combines several things
  - Linux Cgroups
  - Kernel Namespaces
  - Docker Image
  - Has a lifecycle



## Dockerfile

---

- Like a Makefile (shell script with keywords)
- Extends from a Base Image
- Results in a new Docker Image
- Imperative, not Declarative



## Dockerfile Example

---

```
FROM centos:centos6
MAINTAINER aotto@aotto.com
RUN yum -y install openssh-server
EXPOSE 22
ADD start.sh /start.sh
CMD /start.sh
```

## Dockerfile Example

---

```
FROM adrian_server_with_ssh
MAINTAINER aotto@aotto.com
RUN yum -y install httpd
EXPOSE 22 80
ADD start.sh /start.sh
CMD /start.sh
```

# Docker Container Lifecycle

---

- The Life of a Container
  - Conception
    - BUILD an Image from a Dockerfile
  - Birth
    - RUN (create+start) a container
  - Reproduction
    - COMMIT (persist) a container to a new image
    - RUN a new container from an image
  - Sleep
    - KILL a running container
  - Wake
    - START a stopped container
  - Death
    - RM (delete) a stopped container
- Extinction
  - RMI a container image (delete image)

# Docker CLI Commands (v1.4.1)

<b>attach</b>	Attach to a running container	<b>port</b>	Lookup public-facing port
<b>build</b>	<b>Build an image from a Dockerfile</b>	<b>pause</b>	Pause all processes within a container
<b>commit</b>	<b>Create new image from container's changes</b>	<b>ps</b>	<b>List containers</b>
<b>cp</b>	Copy files from containers fs to host	<b>pull</b>	Pull image or repo from docker registry
<b>create</b>	Create a new container	<b>push</b>	Push image or repo to docker registry
<b>diff</b>	Inspect changes on a container's fs	<b>restart</b>	Restart a running container
<b>events</b>	Get real time events from the server	<b>rm</b>	<b>Remove one or more containers</b>
<b>exec</b>	Run a command in a running container	<b>rmi</b>	<b>Remove one or more images</b>
<b>export</b>	Stream contents of container as tar	<b>run</b>	<b>Run a command in a new container</b>
<b>history</b>	Show the history of an image	<b>save</b>	Save an image to a tar archive
<b>images</b>	<b>List images</b>	<b>search</b>	Search for an image in the docker index
<b>import</b>	Create new fs image from a tarball	<b>start</b>	Start a stopped container
<b>info</b>	Display system-wide information	<b>stop</b>	Stop a running container
<b>inspect</b>	Return low-level info on a container	<b>tag</b>	Tag an image into a repository
<b>kill</b>	<b>Kill a running container</b>	<b>top</b>	Lookup running processes of a container
<b>load</b>	Load an image from a tar archive	<b>unpause</b>	Unpause a paused container
<b>login</b>	Login to the docker registry server	<b>version</b>	Show the docker version information
<b>logout</b>	Log out from a Docker registry server	<b>wait</b>	Block and print exit code upon cont exit
<b>logs</b>	Fetch the logs of a container		

## Example Docker Commands

---

- Get an interactive bash shell on a variety of linux distros
  - `docker run -i -t ubuntu:latest bash`
  - `docker run -i -t centos:latest bash`
  - `docker run -i -t debian:latest bash`
- Start a Debian container that runs an Apache server
  - `docker run -d -p 80:80 -p 443:443 httpd:latest`
- Start a Debian container running mysql that uses the `/var/lib/mysql` directory from the host
  - `docker run --name db -p 3306:3306 -e MYSQL_ROOT_PASSWORD=1234.Rack4U2 -v /var/lib/mysql:/var/lib/mysql -d mysql`

# 1.2

## **Lesson One – Part Two**

### **Toolbox Overview**



# Docker Hub

---

- Public Repository for base images and community images
- Private Repositories for sale (get one for free per account)
  - Revenue model for Docker, Inc.
  - Requires authentication to view or update

# Trusted Build System

---

- Docker hosted service
- GitHub source code to docker image repository
- Uses Dockerfile from the GitHub repo

# CoreOS

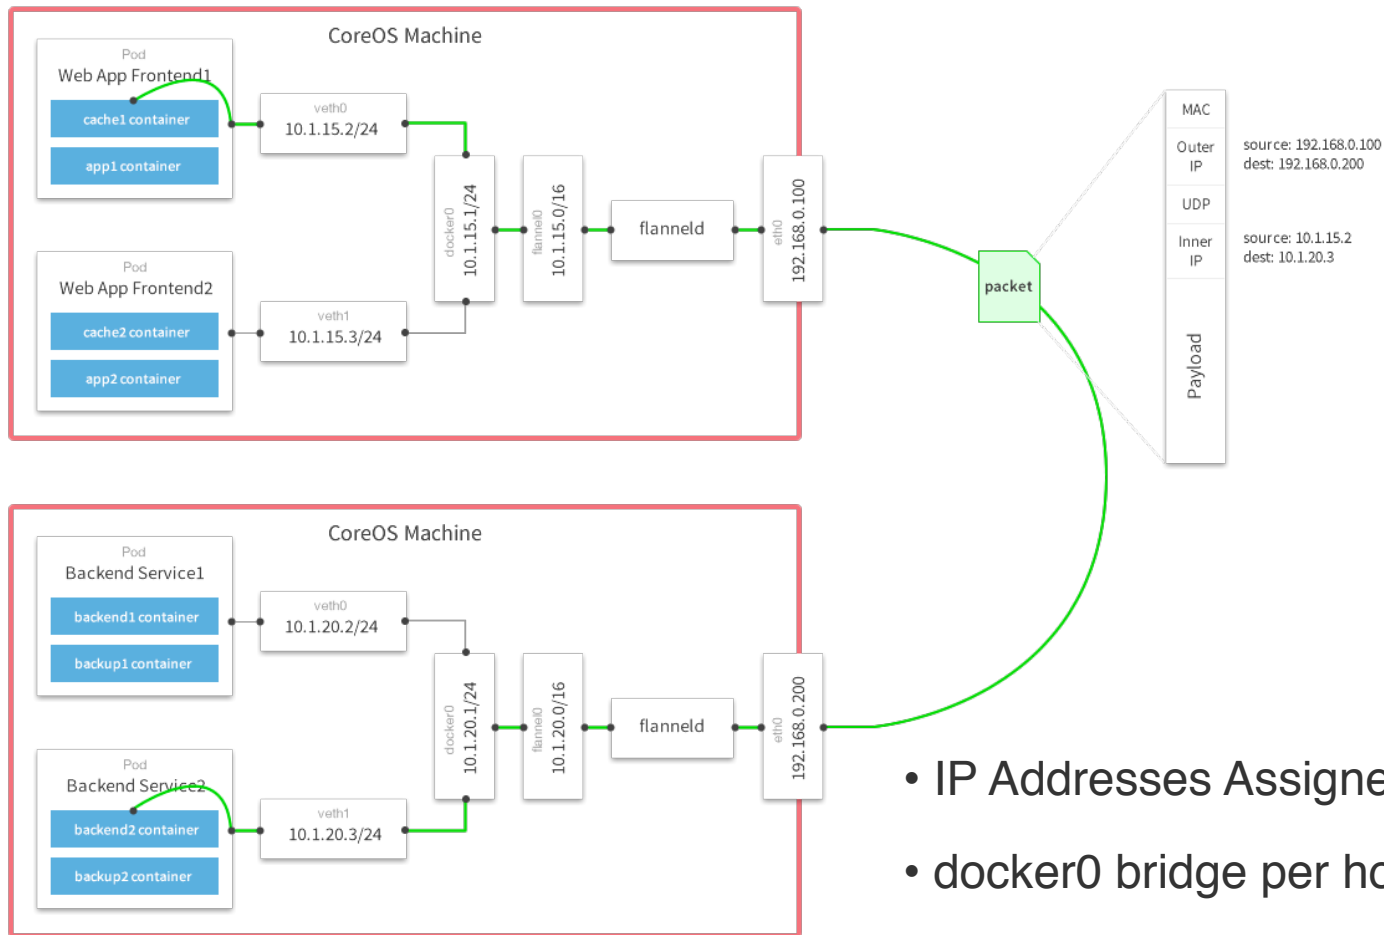
---

- Micro-OS
- Kernel + SSH + Docker + etcd + systemd
- Uses “fleet” to create clusters of CoreOS hosts
- Suitable for running microservices
- Automatic updates that require system level reboots

# Weave and Flannel

---

- Both
  - Overlay network systems
    - Conceptually a way to build a VPN between your containers
  - Build a full mesh container-to-container network using VxLAN or UDP tunnels
- Weave
  - More mature than most of the container networking solutions
  - Expected to work well for deployments of < 200 containers
  - Simple user interface that behaves like “docker run”
- Flannel
  - Part of the CoreOS project, intended for use with Kubernetes
  - Each machine in the cluster is assigned a non-routable subnet.



- IP Addresses Assigned per pod
- docker0 bridge per host

# Kubernetes

---

- Started by Google
- Opinionated Orchestration for Docker Containers
- Declarative expression of the desired end state
- Arranges services in pods
- Appropriate for most container use cases
- Hosts run a “minion” service that reports to the “master” which tracks the cluster state

# How Docker works with OpenStack

---

- Nova-docker
  - Virt driver for Nova
  - Basic start and stop (basic VM lifecycle)
- Heat Resource
  - Create containers on Nova instances
  - Represent them in HOT files
  - No resource scheduler
  - Native management interface only
- Magnum
  - OpenStack Service for Containers as a first class resource
    - Integrates with Keystone
    - Integrates with Heat
    - Multi-Tenant
    - Asynchronous API
  - Concepts for Node, Bay, Pod, Container, Service, Replication Controller

# LAB 1

## Hands-On Lab 1

### Introduction to Docker



## LAB 1.0 – Get Ready

---

- Set your shell environment variables – Region value is case sensitive!

```
export OS_REGION_NAME=IAD
```

```
export OS_USERNAME=jdoe
```

```
export OS_API_KEY=278db780c7d943a1bc0b41ac359d01a8
```

## LAB 1.1

### Do This First

```
export OS_REGION_NAME=IAD
export OS_USERNAME=your_cloud_username
export OS_API_KEY=your_api_key_here
```

- Spin up two new Docker machines on Rackspace Cloud, switch between them.

```
$ docker-machine create -d rackspace machine1
$ docker-machine create -d rackspace machine2
$ docker-machine ls
$ docker-machine active machine1
```

- Get a shell on a Docker host

```
$ docker-machine ssh machine1
```

- Run ad-hoc containers on the fly for experimentation

```
$ docker run --rm -i -t ubuntu:latest bash
$ docker run --rm -i -t centos:latest bash
$ docker run --name one -d ubuntu:latest sleep 1d
$ docker run --name two -d centos:latest sleep 1d
$ docker ps
```

## LAB 1.2 – Get a shell in a container without running sshd

- Get a shell in container “two” without running sshd:

```
$ docker exec -i -t two bash
```

```
$ uname -a
```

← Notice the kernel belongs to the host?

```
$ cat /etc/redhat-release
```

### Alternate Approach - nsenter

```
root@machine1:~# docker run --rm -v /usr/bin:/target
jpetazzo/nsenter
root@machine1:~# cat << "EOF" > enter.sh
PID=`docker inspect --format {{.State.Pid}} $1`
shift
nsenter --mount --uts --ipc --net --pid --target $PID $*
EOF
root@machine1:~# sh -x enter.sh one
```

## LAB 1.3 – Move a container to another host

- Get a hub.docker.com account
  - Browse to <https://hub.docker.com>
  - Click “Sign up with Github” black button (Provide your Github username+password)
  - Create a repo named “private” (Click “Add Repository” -> Repository)
    - Mark it as private, not public
- Move a container from one host to another.

```
$ $(docker-machine env machine1)
$ docker login
$ docker commit one your_username/private:2015-02-26
$ docker images
$ docker push your_username/private:2015-02-26
$ docker kill one; docker rm one; docker rmi your_username/private:2015-02-26
$ $(docker-machine env machine2)
$ docker login
$ docker run --name one -d your_username/private:2015-02-26
$ docker ps
```

Enter your Github username and password

# Break

# 2

## **Lesson Two**

### **Writing Dockerfiles**

## Dockerfile Example (review)

---

```
FROM centos:centos6
MAINTAINER aotto@aotto.com
RUN yum -y install openssh-server
EXPOSE 22
ADD start.sh /start.sh
CMD /start.sh
```

# Dockerfile Directives

---

- **FROM**
  - The `name:tag` of the parent image that this image will be based on
- **MAINTAINER**
  - Optional line for documenting who does the care and feeding for this image. Typically an email address or name.
- **RUN <command>**
  - A command to run in the container to adjust something on this image. This runs at the time of the `docker build` command.
- **EXPOSE**
  - What TCP ports should be exposed on the host when `docker run -P` is used
- **ADD <src>... <dest>**
  - Add files from the `src` directory in the context of the build to the container filesystem at `dest`
- **CMD command param1 param2**
  - Unless otherwise specified by the `docker run` command, run this command when the container is started with this image



## Build an image from a Dockerfile

---

- Create a new directory, and enter it

```
$ mkdir build
```

```
$ cd build
```

- Create a file named Dockerfile in that directory
- Note: Each line of the Dockerfile is executed in a separate intermediate container, followed by an automatic commit
- Build the image

```
$ docker build -t <name>:<tag> .
```

## Demonstration

---

- Copying files into containers using bind mounts
- Copying files out of containers using `docker cp`

demo

## Using Dockerfile and build rather than commit

---

- Commit stores the current state of the container filesystem
- Dockerfiles give us a repeatable way to create a container image

Best Practice: docker build

# LAB 2

## Hands-On Lab 2

### Writing Dockerfiles

## LAB 2.1 – Writing your own Dockerfile

---

- Fork and/or clone the demo repo

```
git clone https://github.com/adrianotto/dockerfile_demo.git
```

- Edit the Dockerfile to change something
  - Maybe add a new file of your own or change what packages are installed

- Run the build.sh script

```
$ sh -x build.sh
```

- Notice the docker build command that runs
- Now add an additional line to the Dockerfile after the existing yum install line:

```
RUN yum -y install vim
```

- Run the build script again
  - Notice that the previous commands are skipped, and cached output is used!

**Break**

# 3

## **Lesson Three**

### **Linking and Networking Containers**

## Understanding Links

---

- Networks namespaces between containers can be bridged together by Docker
- A “Link” is a network path between two containers on the same host
- Syntax:

```
docker run -d --name my_db training/postgres
docker run -d --name web --link my_db:db -p 80:80 training/webapp python app.py
```

- This creates a pair of containers that can connect to each other
- The /etc/hosts file on the ‘web’ container will have an entry for a host named ‘db’
- Processes running inside the ‘web’ container can access the database at the hostname ‘db’



# Understanding Overlay Networks

---

- Overlay networks are like a VPN between containers
- Typical transports are UDP tunneling and VxLAN.
  - Weave uses UDP tunneling
  - Flannel supports both UDP tunneling and VxLAN
- Gives appearance of containers on separate hosts are on the same physical network segment.
- Performance will be reasonable in most cases, but slower than without an overlay network.

## Weave Example <https://github.com/zettio/weave/releases>

```
$ docker-machine ssh machine1
root@machine1:~# cd /usr/bin && wget https://github.com/zettio/weave/releases/download/v0.9.0/weave && chmod +x weave
root@machine1:/usr/bin# weave launch
root@machine1:/usr/bin# weave run 10.0.1.1/24 --name weave1 -d ubuntu sleep 1d
root@machine1:/usr/bin# exit
$ docker-machine ls
NAME          ACTIVE DRIVER      STATE      URL                      SWARM
machine1      *      rackspace   Running    tcp://162.242.239.109:2376
machine2      *      rackspace   Running    tcp://162.242.243.251:2376
$ docker-machine ssh machine2
root@machine2:~# cd /usr/bin; wget https://github.com/zettio/weave/releases/download/v0.9.0/weave && chmod +x weave
root@machine2:/usr/bin# weave launch 162.242.239.109
root@machine2:/usr/bin# weave run 10.0.1.2/24 --name weave2 -d ubuntu sleep 1d
root@machine2:/usr/bin# exit
$ docker exec -i -t weave2 bash
root@587cd615d122:/# ping -c 4 10.0.1.1
PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
64 bytes from 10.0.1.1: icmp_seq=1 ttl=64 time=0.034 ms
64 bytes from 10.0.1.1: icmp_seq=2 ttl=64 time=0.030 ms
64 bytes from 10.0.1.1: icmp_seq=3 ttl=64 time=0.049 ms
64 bytes from 10.0.1.1: icmp_seq=4 ttl=64 time=0.048 ms

--- 10.0.1.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 2997ms
rtt min/avg/max/mdev = 0.030/0.040/0.049/0.009 ms
```

# LAB 3

## Hands-On Lab 3

### Container Linking

## LAB 3.1 – Linking Containers

---

- First, create a container with a database in it:

```
docker run -d --name my_db training/postgres
```

- Now create a second container that connects to the first using a link:

```
docker run -d --name web --link my_db:db -p 80:5000 training/webapp python app.py
```

- Connect to the web container

```
docker exec -i -t web bash
```

```
root@7f4a56aed86d:/opt/webapp# grep db /etc/hosts
```

```
172.17.0.4      db
```

```
root@7f4a56aed86d:/opt/webapp# apt-get install inetutils-ping
```

```
root@7f4a56aed86d:/opt/webapp# ping -c 4 db
```

```
PING db (172.17.0.4): 48 data bytes
```

```
56 bytes from 172.17.0.4: icmp_seq=0 ttl=64 time=0.084 ms
```

```
56 bytes from 172.17.0.4: icmp_seq=1 ttl=64 time=0.087 ms
```

```
56 bytes from 172.17.0.4: icmp_seq=2 ttl=64 time=0.089 ms
```

```
56 bytes from 172.17.0.4: icmp_seq=3 ttl=64 time=0.087 ms
```

```
--- db ping statistics ---
```

```
4 packets transmitted, 4 packets received, 0% packet loss
```

```
round-trip min/avg/max/stddev = 0.084/0.087/0.089/0.000 ms
```

```
root@7f4a56aed86d:/opt/webapp#
```

## LAB 3.2 – Networking Containers

---

```
$ docker-machine ssh machine1
root@machine1:~# cd /usr/bin && wget https://github.com/zettio/weave/releases/download/v0.9.0/weave && chmod +x weave
root@machine1:/usr/bin# weave launch
root@machine1:/usr/bin# weave run 10.0.1.1/24 --name weave1 -d ubuntu sleep 1d
root@machine1:/usr/bin# exit
$ docker-machine ip machine1
162.242.239.109

$ docker-machine ssh machine2
root@machine2:~# cd /usr/bin; wget https://github.com/zettio/weave/releases/download/v0.9.0/weave && chmod +x weave
root@machine2:/usr/bin# weave launch 162.242.239.109
root@machine2:/usr/bin# weave run 10.0.1.2/24 --name weave2 -d ubuntu sleep 1d
root@machine2:/usr/bin# exit
$(docker-machine env machine2)
$ docker exec -i -t weave2 bash
root@587cd615d122:/# ping -c 4 10.0.1.1
PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
64 bytes from 10.0.1.1: icmp_seq=1 ttl=64 time=0.034 ms
64 bytes from 10.0.1.1: icmp_seq=2 ttl=64 time=0.030 ms
64 bytes from 10.0.1.1: icmp_seq=3 ttl=64 time=0.049 ms
64 bytes from 10.0.1.1: icmp_seq=4 ttl=64 time=0.048 ms

--- 10.0.1.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 2997ms
rtt min/avg/max/mdev = 0.030/0.040/0.049/0.009 ms
```

# Summary

---

- LESSON 1 - PART 1 Intro to Docker
  - Docker, Cgroups, Kernel Namespaces, and Docker Images
- LESSON 1 - PART 2 – Toolbox Overview
  - Docker Hub, Trusted build system, Weave, Flannel, etcd, CoreOS, Magnum, Kubernetes, and how Docker works with OpenStack
- Hands-on LAB 1 - Using docker-machine
  - Spin up Docker containers on Rackspace Cloud, Switch between docker hosts.
  - How to run ad-hoc containers on the fly for experimentation, How to get a shell on a Docker host, and a Docker container without running sshd in the container.
  - How to move a container from one host to another.
- LESSON 2 – Writing Dockerfiles
  - Dockerfile Directives
  - Demonstration of copying files into containers using bind mounts, and copying files out of containers using “docker cp”.
  - Why Dockerfiles are better than customized images made manually and stored with commit.
- Hands-on LAB 2 - Writing Dockerfiles.
  - Saving a Dockerfile in a source repo, and using a build script for configuration injection.
- LESSON 3 - Linking and Networking containers
  - Demonstration of how to use links to communicate between containers on the same host. Using weave to communicate between containers on separate hosts.
- Hands on LAB 3 - Linking containers
  - Container linking on the same host, and mesh networking between containers on different hosts.

# Credits

---

- Adrian Otto
- Simon Jakesch
- Thomas Maddox
- Ash Wilson
- Don Schenck

# THANK YOU



**RACKSPACE®** | 1 FANATICAL PLACE, CITY OF WINDCREST | SAN ANTONIO, TX 78218  
**US SALES:** 1-800-961-2888 | **US SUPPORT:** 1-800-961-4454 | **WWW.RACKSPACE.COM**

© RACKSPACE LTD. | RACKSPACE® AND FANATICAL SUPPORT® ARE SERVICE MARKS OF RACKSPACE US, INC. REGISTERED IN THE UNITED STATES AND OTHER COUNTRIES. | [WWW.RACKSPACE.COM](http://WWW.RACKSPACE.COM)